

OmniVec

Universal Vector Ingestion Platform



Turn any data into AI-ready vector search — in minutes, on Azure. No pipeline to build.

WHAT IT OFFERS

- **Minutes, not months** — one command stands up the platform; vectors searchable the same day.
- **Zero glue code** — configure sources, models & destinations; the rest is handled.
- **Works with your stack** — any source, any vector store, your model. No lock-in.
- **Production-grade** — identity, RBAC, autoscaling, metrics — secure from day one.

WHAT YOUR CUSTOMERS CAN BUILD

- **RAG & copilots** — ground LLMs and chatbots in enterprise data.
- **Semantic search** — search by meaning across docs, tickets, wikis.
- **Multi-modal doc search** — PDFs, scans & images via OCR/vision.
- **Always-fresh index** — change feed re-embeds as data changes.
- **Similarity & recs** — find related items, dedupe, cluster.
- **Vectorize what you have** — backfill databases & blobs into an index.

WHAT'S IN THE BOX

- **Universal connectors** — Blob, Cosmos DB, PostgreSQL, MSSQL.
- **Bring-your-own model** — Azure OpenAI or self-hosted GPU.
- **Smart chunking** — truncate or overlap-chunk, per pipeline.
- **Real-time sync** — change feed & Event Grid keep vectors current.
- **Three interfaces** — Web UI, CLI and REST API.
- **Elastic & cost-aware** — autoscaling; GPU scale-to-zero.

WHY OMNIVEC

BUILD IT YOURSELF

- ✗ Weeks of glue code
- ✗ Connector per source
- ✗ Manual re-indexing
- ✗ DIY GPU & scaling
- ✗ You own security
- ✗ Hard to demo early

WITH OMNIVEC

- ✓ Deploy in ~20 min
- ✓ Universal connectors
- ✓ Real-time auto-sync
- ✓ Managed GPUs
- ✓ Secure by default
- ✓ Demo the same day

WORKS WITH YOUR DATA

SOURCES

EMBED

VECTORS

Blob · Cosmos · Postgres · SQL → AOAI / GPU → Cosmos · pgvector · SQL

- **One-command deploy** — azd up provisions AKS, Cosmos DB, Key Vault & more (~15–25 min).
- **Multi-modal** — text, images and PDFs (vision + OCR), truncate or chunk.
- **Data stays yours** — self-host models so content never leaves your tenant.
- **Observable & elastic** — health, per-pipeline metrics; GPU scale-to-zero.

ROADMAP

- **Hardening to GA** — stable APIs, fuller feature set.
- **Network isolation** — private endpoints + WAF.
- **Zero-trust cluster** — mTLS / mesh + cosign signing.
- **More connectors** — broader source & dest. coverage.

WHY IT WINS

~20

min to demo

customer POC, end to end

\$5–10

per day

default footprint, no GPU

0

lines of glue

configuration, not code

3

ways to drive

Web UI · CLI · REST API